

## Chapter 1

# Introduction

Parser combinators [Hutton 1992] are an elegant approach for writing parsers in a manner that remains close to their original grammar specification. `parsley` [Willis and Wu 2018] is a parser combinator library implemented as an embedded domain-specific language (DSL) [Hudak 1996] in Scala, with an API inspired by the `parsec` [Leijen and Meijer 2001] family of libraries in Haskell. However, as with many libraries, there exists a learning curve to utilising `parsley` and parser combinator libraries in an idiomatic manner.

While well-documented, the wealth of information to get started with `parsley` can be overwhelming for users, particularly those new to parser combinators. Although `parsley` itself has a user-friendly API, parser combinators in general have pitfalls that may be unexpected for new users. Even experienced users can unintentionally write unidiomatic parsers: `parsley` has first-class support for a number of design patterns [Willis and Wu 2022] for writing maintainable parsers, which users migrating from other parser combinator libraries may not be aware of.

This project aims to address these issues by developing a companion *linting* tool for `parsley`, called `parsley-garnish`, that provides automated code hints and fixes to assist users in writing idiomatic and correct parsers. A number of modern integrated development environments (IDEs) provide code hints to warn programmers about problems in their source code, highlighting offending snippets and suggesting actions to improve suboptimal or incorrect code [Kurbatova et al. 2021]. Many of these linters are designed to detect general issues for the host language, rather than specifically for libraries. However, tools may also utilise domain-specific code analyses in order to detect issues specific to a particular system or problem domain [Renggli et al. 2010; Gregor and Schupp 2006]. Well-designed linters can offer significant benefits to users:

- Linters can be particularly valuable for uncovering subtle issues that might be hard to diagnose and locate, especially in large codebases. Automated fixes can save further effort by resolving issues without manual intervention.
- Linters are also beneficial for teaching best practices in context, offering relevant hints and improvements precisely where sub-optimal code is detected.

For example, suppose a user wants to write a simple arithmetic expression parser in `parsley`, which evaluates the parsed expression as a floating-point calculation. The parser will be based on the following EBNF grammar, with standard arithmetic operator precedence and left-associativity:

$$\begin{aligned} \langle \text{digit} \rangle &::= '0' \dots '9' \\ \langle \text{number} \rangle &::= \langle \text{digit} \rangle + \\ \langle \text{expr} \rangle &::= \langle \text{expr} \rangle '+' \langle \text{term} \rangle \mid \langle \text{expr} \rangle '-' \langle \text{term} \rangle \mid \langle \text{term} \rangle \\ \langle \text{term} \rangle &::= \langle \text{term} \rangle '*' \langle \text{atom} \rangle \mid \langle \text{term} \rangle '/' \langle \text{atom} \rangle \mid \langle \text{atom} \rangle \\ \langle \text{atom} \rangle &::= '(' \langle \text{expr} \rangle ')' \mid \langle \text{number} \rangle \end{aligned}$$

By closely following the structure of the grammar, a naïve first attempt at writing the parser-evaluator in `parsley` may resemble the following:

```
val number: Parsley[Float] = digit.foldLeft1(0)((n, d) => n * 10 + d.asDigit).map(_.toFloat)

lazy val expr: Parsley[Float] = (expr, char('+') ~> term).zipped(_ + _)
                                | (expr, char('-') ~> term).zipped(_ - _)
                                | term
lazy val term: Parsley[Float] = (term, char('*') ~> atom).zipped(_ * _)
                                | (term, char('/') ~> atom).zipped(_ / _)
                                | atom
lazy val atom: Parsley[Float] = char('(') ~> expr <~ char(')') | number
```



## Outline

?? begins this report by outlining the key background material required to understand the project. This includes an introduction to linters and how they are typically implemented; a discussion on writing linters specifically for Scala; and an overview of parser combinators, their design patterns, and the parsley library.

From there, ?? dives into the first set of linting rules that parsley-garnish implements. These relatively simple rules focus on enforcing idiomatic design patterns relating to *implicit conversions* in Scala, and were largely inspired by common issues I've seen in my experiences as a teaching assistant for the second-year undergraduate Wacc compilers project at Imperial.

?? tackles a more ambitious rule to automatically refactor left-recursive parsers into a form that parsley can handle. This chapter introduces the idea of utilising intermediate AST representations to perform higher-level domain-specific transformations. It shows that a basic implementation is not sufficient to handle the complexities of the problem, and motivates the need for a more sophisticated approach. ?? therefore further develops and refines these ideas, drawing inspiration from metaprogramming techniques and parsley's own optimisation machinery.

These improvements allow the left-recursion transformation to be revisited and completed in ??, which also explores more advanced rules that are unlocked by the new AST. Finally, ?? evaluates **TODO**. The remainder of the report concludes with discussion of related work and future directions for the project.